

REINFORCEMENT LEARNING Exercise 4



This week, we use the methods from last week for control. You find tests for the following tasks in `exercise-04_test.py`. Additionally, there is a visualization script. You again need matplotlib. Run the tests and the visualization by the usual commands.

0 Lecture

Watch *Lecture 05: Value Function Approximation*¹ before the upcoming session on Friday, November 23.

1 Off-Policy MC Control with Importance Sampling

This task is again based on the Blackjack environment from the lecture and the last exercise. Implement First Visit Off-Policy MC Control with Importance Sampling,

```
mc_control_importance_sampling(env, num_episodes, behavior_policy, discount_factor=1.0),  
in off_policy_mc.py.
```

2 Q-learning

The tests for this task are based on the Cliff Walking example from the lecture. An implementation can be found in `cliff_walking.py`. Implement the Q-learning algorithm,

```
q_learning(env, num_episodes, discount_factor=1.0, alpha=0.5, epsilon=0.1),  
in q_learning.py.
```

3 Experiences

Make a post in thread *Week 04: Model-free Control* in the forum², where you provide a brief summary of your experience with this exercise, the corresponding lecture and the last meeting.

¹https://ilias.uni-freiburg.de/goto.php?target=xvid_1121350&client_id=unifreiburg

²https://ilias.uni-freiburg.de/goto.php?target=frm_1121060&client_id=unifreiburg