

REINFORCEMENT LEARNING Exercise 7



0 Lecture

Watch *Lecture 09: Exploration and Exploitation*¹ before the upcoming session on Friday, January 11.

1 PPO

Algorithm 1 Proximal Policy Optimization

```

1: procedure PPO
2:   initialize parameters  $\theta^\pi$  of policy  $\pi$  arbitrarily
3:   initialize parameters  $\theta^V$  of value function  $V$  arbitrarily
4:   update parameters  $\theta^{\pi_{\text{old}}}$  of policy  $\pi_{\text{old}}$ 
5:   for iteration=1,2,... do
6:     Run policy  $\pi_{\text{old}}$  for  $e$  episodes
7:     Compute advantage estimates  $\hat{A}_i$ 
8:     for some epochs do
9:       train  $V$ 
10:      Optimize surrogate  $L^{\text{CLIP}}(\theta^\pi) = \hat{\mathbf{E}}_t[\min(\frac{\pi(a_t|s_t)}{\pi_{\text{old}}(a_t|s_t)} \hat{A}_i, \text{clip}(\frac{\pi(a_t|s_t)}{\pi_{\text{old}}(a_t|s_t)}, 1 - \epsilon, 1 + \epsilon) \hat{A}_i)]$ 
11:    update parameters  $\theta^{\pi_{\text{old}}}$  of policy  $\pi_{\text{old}}$ 

```

Implement the `ppo` function in `ppo.py` using the *Clipped Surrogate Objective*, where `update_frequency` is the amount of episodes between updates and `epochs` is the number of full iterations over the collected data, which can be implemented using the `next_batch(index, batch_size)` method of the `ReplayBuffer` class. Evaluate again on the modified `MountainCar` environment.

2 Experiences

Make a post in thread *Week 08: Advanced Policy Gradient Algorithms* in the forum², where you provide a brief summary of your experience with this exercise, the corresponding lecture and the last meeting. We wish you a nice Winter break and a happy new year.

¹ https://ilias.uni-freiburg.de/goto.php?target=xvid_1121354&client_id=unifreiburg

² https://ilias.uni-freiburg.de/goto.php?target=frm_1121060&client_id=unifreiburg